

Data Visualization and SVM BY selecting random features

```
# Import necessary libraries
```

```
import pandas as pd
```

```
import matplotlib.pyplot as plt
```

```
import seaborn as sns
```

```
from sklearn.model_selection import train_test_split
```

```
from sklearn.svm import SVC
```

```
from sklearn.preprocessing import LabelEncoder, StandardScaler
```

```
from sklearn.metrics import classification_report, confusion_matrix
```

```
import numpy as np
```

```
from sklearn.metrics import accuracy_score
```

```
# Specify the file path
```

```
file_path = '/content/drive/MyDrive/train.csv'
```

```
# Load the dataset
```

```
data = pd.read_csv(file_path)
```

```
print(data.head)
```

```
# Specify the columns to extract (replace with actual column names)
```

```
columns_to_extract = ['LotArea', 'MSZoning', 'SalePrice'] # Replace with your desired  
column names
```

```
# Select the specific columns and first 20 rows
```

```
subset_data = data[columns_to_extract].head(100)
```

```
# Display the data as a table
print("Subset of the data (3 columns, 100 rows):")
display(subset_data)

# Visualize the data with plots
# 1. Bar Plot for one of the columns
plt.figure(figsize=(10, 6))
sns.barplot(data=subset_data, x='LotArea', y='SalePrice') # Replace with actual column
names
plt.title('Bar Plot of AreaLot vs SalePrice')
plt.xlabel('LotArea')
plt.ylabel('SalePrice')
plt.xticks(rotation=45)
plt.show()

# 2. Scatter Plot for two columns
plt.figure(figsize=(10, 6))
sns.scatterplot(data=subset_data, x='LotArea', y='SalePrice', hue='MSZoning',
palette='viridis') # Replace as needed
plt.title('Scatter Plot of Column1 vs Column3')
plt.xlabel('LotArea')
plt.ylabel('SalePrice')
plt.show()

# Select only numeric columns for the heatmap
numeric_data = subset_data.select_dtypes(include=['float64', 'int64'])

# Check if there are numeric columns
if not numeric_data.empty:
    # Create the heatmap
```

```

plt.figure(figsize=(8, 5))

sns.heatmap(numeric_data.corr(), annot=True, cmap='coolwarm', fmt='.2f')

plt.title('Correlation Heatmap of Selected Columns')

plt.show()

else:

    print("No numeric columns available for heatmap.")


# Select numeric columns for features and one column for the target
columns_to_extract = ['GrLivArea', 'TotalBsmtSF', 'SalePrice'] # Update with other column
names

subset_data = data[columns_to_extract].dropna().head(100) # Ensure no missing values and
use first 100 rows


# Prepare features (GrLivArea, TotalBsmtSF) and target (SalePrice)
features = subset_data[['GrLivArea', 'TotalBsmtSF']] # Two numeric columns
target = subset_data['SalePrice']


# Bin SalePrice into categories: 'Low', 'High'
subset_data['PriceCategory'] = pd.cut(target, bins=2, labels=['Low', 'High'])
target_encoded = LabelEncoder().fit_transform(subset_data['PriceCategory'])


# Standardize numeric features
scaler = StandardScaler()

features_scaled = scaler.fit_transform(features)


# Train-test split
X_train, X_test, y_train, y_test = train_test_split(features_scaled, target_encoded,
test_size=0.3, random_state=42)


# Create and train Linear SVM model

```

```

svm_model = SVC(kernel='linear', random_state=42)
svm_model.fit(X_train, y_train)

# Make predictions
y_pred = svm_model.predict(X_test)

# Evaluate the model
print("Classification Report:")
print(classification_report(y_test, y_pred))

print("\nConfusion Matrix:")
print(confusion_matrix(y_test, y_pred))

# Visualize the linear hyperplane
plt.figure(figsize=(10, 6))

# Scatter plot of features with their labels
for i, label in enumerate(['Low', 'High']):
    plt.scatter(features_scaled[target_encoded == i, 0], features_scaled[target_encoded == i,
1], label=label)

# Plot the decision boundary
coef = svm_model.coef_[0]
intercept = svm_model.intercept_[0]
slope = -coef[0] / coef[1]
xx = np.linspace(features_scaled[:, 0].min(), features_scaled[:, 0].max())
yy = slope * xx - intercept / coef[1]
plt.plot(xx, yy, 'k--', label='Hyperplane')
coef = svm_model.coef_[0]
intercept = svm_model.intercept_[0]

```

```

# Define slope and intercept of the hyperplane
slope = -coef[0] / coef[1]
threshold_intercept = -intercept / coef[1]

print(f"Slope of Hyperplane: {slope}")
print(f"Threshold Intercept on Hyperplane: {threshold_intercept}")
plt.title("Linear SVM Decision Boundary")
plt.xlabel("GrLivArea (scaled)")
plt.ylabel("TotalBsmtSF (scaled)")
plt.legend()
plt.grid()
plt.show()

x1= subset_data['GrLivArea']
y = subset_data['SalePrice']

# Calculate correlation coefficient
r = np.corrcoef(x1, y)[0, 1]
print(f"Correlation Coefficient (GrLivArea vs SalePrice): {r:.2f}")

x2 = subset_data['TotalBsmtSF']
y = subset_data['SalePrice']
r = np.corrcoef(x2, y)[0, 1]
print(f"Correlation Coefficient (TotalBsmtSF vs SalePrice): {r:.2f}")

# Select only numeric columns
numeric_data = data.select_dtypes(include=['float64', 'int64'])

```

```

if 'SalePrice' in numeric_data.columns:

    # Calculate the correlation of each column with SalePrice
    saleprice_correlation = numeric_data.corr()['SalePrice'].sort_values(ascending=False)

    # Display the correlation values
    print("Correlation of each feature with SalePrice:")
    print(saleprice_correlation)
else:
    print("Column 'SalePrice' not found in the dataset.")

accuracy = accuracy_score(y_test, y_pred)
print(f"Model Accuracy: {accuracy:.2f}")

```

Output :

```

<bound method NDFrame.head of      Id MSSubClass MSZoning LotFrontage LotArea
Street Alley LotShape \
0      1      60      RL      65.0   8450  Pave  NaN      Reg
1      2      20      RL      80.0   9600  Pave  NaN      Reg
2      3      60      RL      68.0  11250  Pave  NaN      IR1
3      4      70      RL      60.0   9550  Pave  NaN      IR1
4      5      60      RL      84.0  14260  Pave  NaN      IR1
...  ...      ...      ...      ...      ...      ...      ...
1455 1456      60      RL      62.0   7917  Pave  NaN      Reg
1456 1457      20      RL      85.0  13175  Pave  NaN      Reg
1457 1458      70      RL      66.0   9042  Pave  NaN      Reg
1458 1459      20      RL      68.0   9717  Pave  NaN      Reg
1459 1460      20      RL      75.0   9937  Pave  NaN      Reg

LandContour Utilities ... PoolArea PoolQC Fence MiscFeature MiscVal \
0      Lvl  AllPub ...      0  NaN  NaN      NaN      0

```

1	Lvl	AllPub	...	0	NaN	NaN	NaN	0
2	Lvl	AllPub	...	0	NaN	NaN	NaN	0
3	Lvl	AllPub	...	0	NaN	NaN	NaN	0
4	Lvl	AllPub	...	0	NaN	NaN	NaN	0
...
1455	Lvl	AllPub	...	0	NaN	NaN	NaN	0
1456	Lvl	AllPub	...	0	NaN	MnPrv	NaN	0
1457	Lvl	AllPub	...	0	NaN	GdPrv	Shed	2500
1458	Lvl	AllPub	...	0	NaN	NaN	NaN	0
1459	Lvl	AllPub	...	0	NaN	NaN	NaN	0

	MoSold	YrSold	SaleType	SaleCondition	SalePrice
0	2	2008	WD	Normal	208500
1	5	2007	WD	Normal	181500
2	9	2008	WD	Normal	223500
3	2	2006	WD	Abnorml	140000
4	12	2008	WD	Normal	250000
...
1455	8	2007	WD	Normal	175000
1456	2	2010	WD	Normal	210000
1457	5	2010	WD	Normal	266500
1458	4	2010	WD	Normal	142125
1459	6	2008	WD	Normal	147500


[1460 rows x 81 columns]>


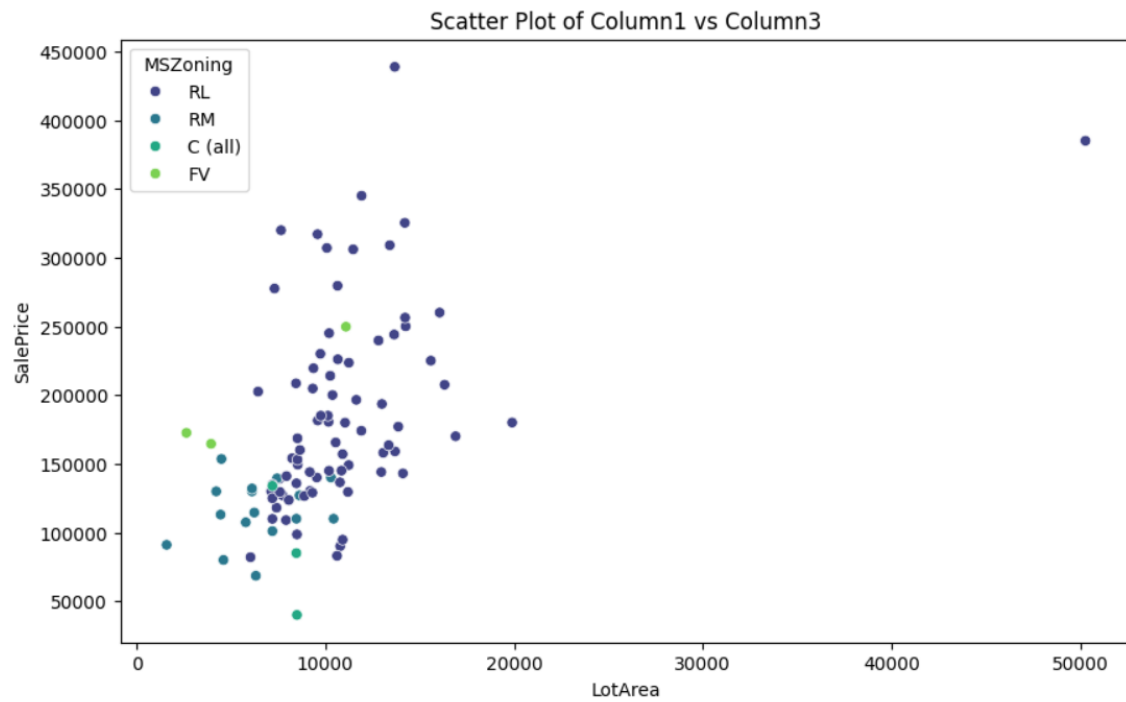
Subset of the data (3 columns, 100 rows):

	LotArea	MSZoning	SalePrice
0	8450	RL	208500

	LotArea	MSZoning	SalePrice
1	9600	RL	181500
2	11250	RL	223500
3	9550	RL	140000
4	14260	RL	250000
...
95	9765	RL	185000
96	10264	RL	214000
97	10921	RL	94750
98	10625	RL	83000
99	9320	RL	128950

100 rows × 3 columns






Classification Report:

precision recall f1-score support

0	0.50	1.00	0.67	4
1	1.00	0.85	0.92	26

accuracy		0.87		30
macro avg	0.75	0.92	0.79	30
weighted avg	0.93	0.87	0.88	30



Confusion Matrix:

[[4 0]

[4 22]]

Slope of Hyperplane: -0.9036829820979518

Threshold Intercept on Hyperplane: 1.2492900768016308

Correlation Coefficient of each column with output to get better Accuracy

Correlation Coefficient (GrLivArea vs SalePrice): 0.74

Correlation Coefficient (TotalBsmtSF vs SalePrice): 0.62

Correlation of each feature with SalePrice:

SalePrice	1.000000
OverallQual	0.790982
GrLivArea	0.708624
GarageCars	0.640409
GarageArea	0.623431
TotalBsmtSF	0.613581
1stFlrSF	0.605852
FullBath	0.560664
TotRmsAbvGrd	0.533723
YearBuilt	0.522897
YearRemodAdd	0.507101
GarageYrBlt	0.486362
MasVnrArea	0.477493
Fireplaces	0.466929
BsmtFinSF1	0.386420
LotFrontage	0.351799
WoodDeckSF	0.324413
2ndFlrSF	0.319334
OpenPorchSF	0.315856
HalfBath	0.284108
LotArea	0.263843
BsmtFullBath	0.227122
BsmtUnfSF	0.214479

BedroomAbvGr 0.168213
ScreenPorch 0.111447
PoolArea 0.092404
MoSold 0.046432
3SsnPorch 0.044584
BsmtFinSF2 -0.011378
BsmtHalfBath -0.016844
MiscVal -0.021190
Id -0.021917
LowQualFinSF -0.025606
YrSold -0.028923
OverallCond -0.077856
MSSubClass -0.084284
EnclosedPorch -0.128578
KitchenAbvGr -0.135907
Name: SalePrice, dtype: float64
Model Accuracy: 0.87

Highest Accuracy By choosing the most correlated columns as input

```
# Select numeric columns for features and one column for the target with highest accuracy
columns_to_extract = ['OverallQual', 'GrLivArea', 'SalePrice'] # Update with other column names

subset_data = data[columns_to_extract].dropna().head(100) # Ensure no missing values and use first 100 rows

# Prepare features (GrLivArea, TotalBsmtSF) and target (SalePrice)
```

```
features = subset_data[['OverallQual','GrLivArea']] # Two numeric columns
target = subset_data['SalePrice']

# Bin SalePrice into categories: 'Low', 'High'
subset_data['PriceCategory'] = pd.cut(target, bins=2, labels=['Low', 'High'])
target_encoded = LabelEncoder().fit_transform(subset_data['PriceCategory'])

# Standardize numeric features
scaler = StandardScaler()
features_scaled = scaler.fit_transform(features)

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(features_scaled, target_encoded,
test_size=0.3, random_state=42)

# Create and train Linear SVM model
svm_model = SVC(kernel='linear', random_state=42)
svm_model.fit(X_train, y_train)

# Make predictions
y_pred = svm_model.predict(X_test)

# Evaluate the model
print("Classification Report:")
print(classification_report(y_test, y_pred))

print("\nConfusion Matrix:")
print(confusion_matrix(y_test, y_pred))

# Visualize the linear hyperplane
```

```

plt.figure(figsize=(10, 6))

# Scatter plot of features with their labels
for i, label in enumerate(['Low', 'High']):
    plt.scatter(features_scaled[target_encoded == i, 0], features_scaled[target_encoded == i,
1], label=label)

# Plot the decision boundary
coef = svm_model.coef_[0]
intercept = svm_model.intercept_[0]
slope = -coef[0] / coef[1]
xx = np.linspace(features_scaled[:, 0].min(), features_scaled[:, 0].max())
yy = slope * xx - intercept / coef[1]
plt.plot(xx, yy, 'k--', label='Hyperplane')
coef = svm_model.coef_[0]
intercept = svm_model.intercept_[0]

# Define slope and intercept of the hyperplane
slope = -coef[0] / coef[1]
threshold_intercept = -intercept / coef[1]

print(f"Slope of Hyperplane: {slope}")
print(f"Threshold Intercept on Hyperplane: {threshold_intercept}")
plt.title("Linear SVM Decision Boundary")
plt.xlabel("GrLivArea (scaled)")
plt.ylabel("OverallQual (scaled)")
plt.legend()
plt.grid()
plt.show()
accuracy = accuracy_score(y_test, y_pred)

```

```
print(f"Model Accuracy: {accuracy:.2f}")
```

Output :

Classification Report:

	precision	recall	f1-score	support
0	0.75	0.75	0.75	4
1	0.96	0.96	0.96	26
accuracy			0.93	30
macro avg	0.86	0.86	0.86	30
weighted avg	0.93	0.93	0.93	30

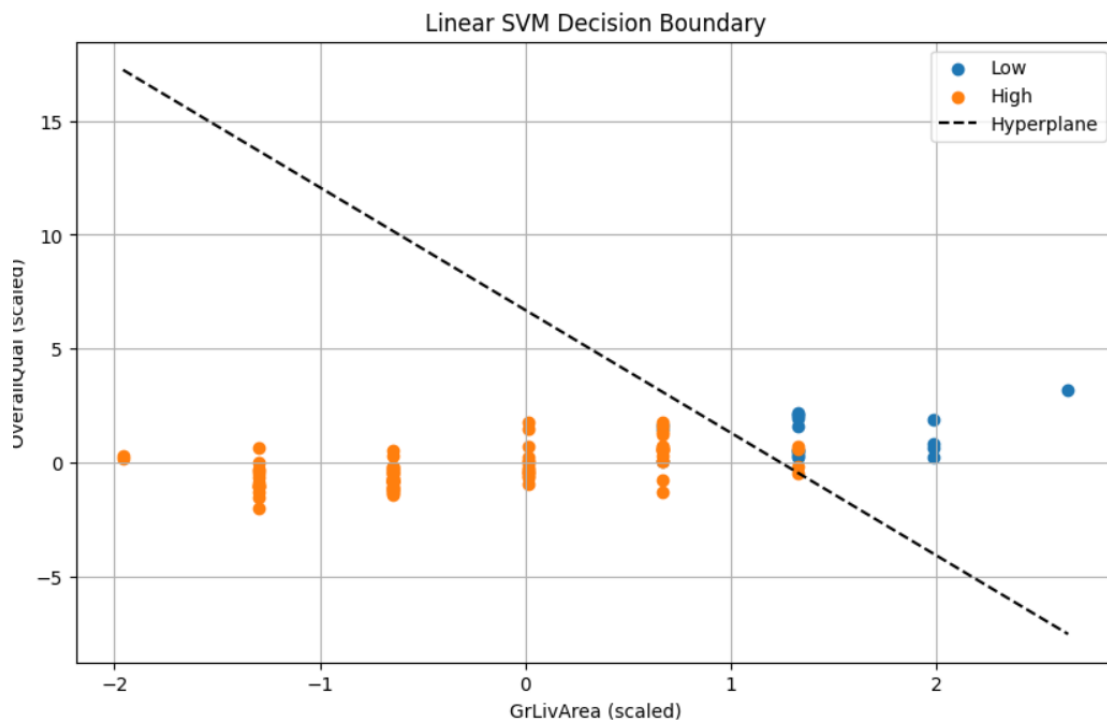
Confusion Matrix:

```
[[ 3  1]
```

```
 [ 1 25]]
```

Slope of Hyperplane: -5.383770452884379

Threshold Intercept on Hyperplane: 6.6949558734721775



Model Accuracy: 0.93

Output when taking less correlated inputs

Select numeric columns for features and one column for the target

```
columns_to_extract = ['LotFrontage', 'BsmtFinSF1', 'SalePrice'] # Update with other column names
```

```
subset_data = data[columns_to_extract].dropna().head(100) # Ensure no missing values and use first 100 rows
```

Prepare features (GrLivArea, TotalBsmtSF) and target (SalePrice)

```
features = subset_data[['LotFrontage', 'BsmtFinSF1']] # Two numeric columns
```

```
target = subset_data['SalePrice']
```

Bin SalePrice into categories: 'Low', 'High'

```
subset_data['PriceCategory'] = pd.cut(target, bins=2, labels=['Low', 'High'])
```

```
target_encoded = LabelEncoder().fit_transform(subset_data['PriceCategory'])
```



```
# Standardize numeric features

scaler = StandardScaler()

features_scaled = scaler.fit_transform(features)


# Train-test split

X_train, X_test, y_train, y_test = train_test_split(features_scaled, target_encoded,
test_size=0.3, random_state=42)


# Create and train Linear SVM model

svm_model = SVC(kernel='linear', random_state=42)

svm_model.fit(X_train, y_train)


# Make predictions

y_pred = svm_model.predict(X_test)


# Evaluate the model

print("Classification Report:")

print(classification_report(y_test, y_pred))


print("\nConfusion Matrix:")

print(confusion_matrix(y_test, y_pred))


# Visualize the linear hyperplane

plt.figure(figsize=(10, 6))


# Scatter plot of features with their labels

for i, label in enumerate(['Low', 'High']):

    plt.scatter(features_scaled[target_encoded == i, 0], features_scaled[target_encoded == i,
1], label=label)
```

```

# Plot the decision boundary
coef = svm_model.coef_[0]
intercept = svm_model.intercept_[0]
slope = -coef[0] / coef[1]
xx = np.linspace(features_scaled[:, 0].min(), features_scaled[:, 0].max())
yy = slope * xx - intercept / coef[1]
plt.plot(xx, yy, 'k--', label='Hyperplane')
coef = svm_model.coef_[0]
intercept = svm_model.intercept_[0]

# Define slope and intercept of the hyperplane
slope = -coef[0] / coef[1]
threshold_intercept = -intercept / coef[1]

print(f"Slope of Hyperplane: {slope}")
print(f"Threshold Intercept on Hyperplane: {threshold_intercept}")
plt.title("Linear SVM Decision Boundary")
plt.xlabel("LotFrontage (scaled)")
plt.ylabel("BsmtFinSF1 (scaled)")
plt.legend()
plt.grid()
plt.show()
accuracy = accuracy_score(y_test, y_pred)
print(f"Model Accuracy: {accuracy:.2f}")

```

Output :

Classification Report:

	precision	recall	f1-score	support
0	0.00	0.00	0.00	9

1 0.70 1.00 0.82 21

accuracy 0.70 30

macro avg 0.35 0.50 0.41 30

weighted avg 0.49 0.70 0.58 30

Confusion Matrix:

```
[[ 0  9]
```

```
 [ 0 21]]
```

Slope of Hyperplane: 0.2311629509675174

Threshold Intercept on Hyperplane: -4753.151100667716

/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565:

UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

```
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```

/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565:

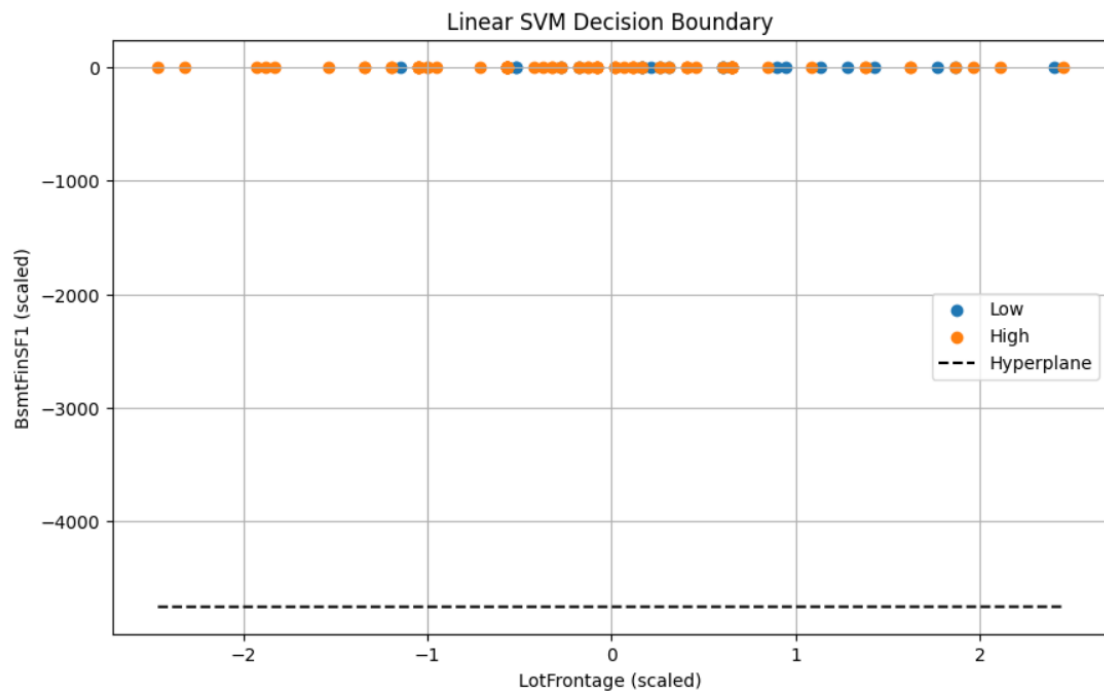
UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

```
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```

/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565:

UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

```
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```



Model Accuracy: 0.70